

Userspace Drivers



ECE 373

What is a userspace driver?

What is a userspace driver?

- A way to drive hardware from outside kernel
- User space program
- Accesses various resources
 - I/O ports
 - Memory regions
 - Control interfaces
- Resources presented by kernel scanning hardware in standardized ways
 - e.g. PCI bus scan
- Not a .ko file!

What isn't a userspace driver?

- Scripts/programs using only sysfs or proc
- Scripts/programs using devfs
 - If /dev entry exists, likely a kernel driver is present
- Exceptions do exist though
 - Graphics drivers can be combination of user and kernel mode
 - Nvidia kmod-nvidia plus Xorg userspace driver
 - More on this later

Drivers of the rings

- Kernel runs in ring 0 on CPU
- Hardware is protected by ring 0 access
- Userspace runs in ring 3 on CPU
- Userspace drivers run in userspace
- But system calls...



What's not available?

- Direct access to all kernel memory
 - Not always a bad thing
 - Requires more maps if you need access
- Interrupts
 - Only fire in kernel space
 - Priority is lost in userspace
- Schedule priority can be a problem
 - Just “another” userspace process

But first...



What is mapping memory?

- First, man `mmap()`
- Maps files into memory...obviously
- Access with array/pointer operations
- Why is this useful?
- Why not just use `open()`, `read()`, `write()`?
- What's the tradeoff?



Files in memory


- Advantage to access large files
 - Reading from a file using memory operations can be quick
- Kernel can manage page resources better
 - mmap'd pages backed by page cache
 - Can swap pieces out quickly, versus file page cache writing back before swapping
- Disadvantage – uses up memory map space
 - Limited space in 32-bit addressing

Files in memory, cont.

- Great for sharing file data between processes
 - MAP_SHARED, MAP_ANONYMOUS attributes
 - Page cached once, rewritten on updates, everyone sees update
- Think web server...
 - Multiple threads on multiple files
- Other examples?


What's it look like?

- What a normal file looks like when opened



What's it look like?

- When a file is mmap'd



Mapping a file is as simple as...

- `open()` the file
- `mmap()` the file
- **Inspect** the memory region returned from `mmap()`
- `munmap()` the file
- `close()` the file
- **Example!**

Why use a userspace driver?

Why use a userspace driver?

- Quick prototyping
- Free from kernel ABI changes
 - Application Binary Interface
- Typically doesn't blow up the machine
 - Constrained to mapped space
- Ease of use
 - Can be written in most compiled and scripted languages
- May not have ability to change the kernel

Why not to use a userspace driver?

Why not to use a userspace driver?

- Performance and latency on wake-ups (more on this in a bit...)
- Power efficiency
- No interrupts
- No direct access to kernel infrastructure
 - Kernel times, jiffies, physical memory, traditional DMA, etc.

Userspace drivers in embedded

- Popular for fully-encapsulated devices
 - Touchscreen
 - Battery driver maybe?
- Popular for simple devices
 - LEDs, GPIOs, etc.
 - Arduino, BeagleBone, etc
- Not popular when interfacing with subsystems
 - Networking, audio
- DPDK – Data Plane Development Kit
 - Intel SDK to enable switching infrastructure, kernel bypass

The hybrid approach

- UIO driver a good example
 - <https://elixir.bootlin.com/linux/latest/source/drivers/uio>
- RDMA drivers
 - OFED stack
- Graphics a perfect example
 - Lightweight portion in kernel, hides some HW interfaces, power management, etc.
 - Heavyweight portion in userspace, interfacing with X server

Performance considerations

- Pure userspace drivers can be quick
 - No interrupts means...?
- Polling can have negative side effects (hurry up and wait!)
- Polling very good for latency-sensitive operations
 - Good for stock trading
 - Good for stop-lights working properly
 - Good for car computers getting the “brake pressed” signal



What about the kernel?

- Pure userspace approach
 - Must break into the kernel somehow
 - Kernel provides mechanism to get at memory
- Hybrid approach
 - Small shim driver in kernel space
 - Exports driver interface file (/dev) for userspace to drive

Mapping into the kernel

- `/dev/mem`
- It's a file, right?
- `man mem...`
- `open ()` and `mmap ()`
- But where to map?
- May require some pleading with the kernel...
 - Linux boot parameter: `iomem=relaxed`

Mapping sensitive areas

- strace
- dmidecode
- SMBIOS
 - http://en.wikipedia.org/wiki/System_Management_BIOS
- Follow dmidecode...

SecureBoot, oh my

- Attempt at securing EFI to bootloader security
- Trying to prevent rootkit access to the bootloader and OS
- “Signed” modules used, signed by a CA
 - Guess who's the CA...
- Required for Win8/10 certification on ARM only
- Red Hat, SuSE, others still moving forward with SecureBoot compliance on x86

What does this mean to me?

- SecureBoot wants to lock access to kernel space down
- /dev/mem accesses likely to be prohibited
- Kernel modules (e.g. drivers) will need to be signed
- Details still being worked

Let's mess with hardware

- To `mmap()` successfully on `/dev/mem`, need a place to map
- What are we trying to map on our PCI device to control it?
 - Hint: what do we map in the kernel?
- Where do we find the BAR address?
 - `Lspci...`
 - Can look in `/sysfs...`

The basics...

- Great examples exist out there
- E1000regs, aka ethregs, scans bus for devices
- Find device we want
- Grab BAR address from struct, and away we go
- Basic example

What about interrupts??

- Poll, poll, poll
 - Maybe assign one core in a multi-core machine for this
- Using hybrid approach, interrupt routine could only consist of waking userspace thread
- Can generate interrupts if underlying HW supports it
 - Writing register to fire SW-originated interrupt
- New Sapphire Rapids (Intel) CPUs have userspace-aimed interrupts...TBD how well they work!

Reading and an assignment!

- LDD3 pages 37-39
- Essential pages 558-564
- Man pages on `mmap()`
- Read the source!
 - e1000regs (aka ethregs) on e1000.sourceforge.net
- Assignment #5